

Corporate AI (ChatGPT) Onboarding Training Package

Format: 2 sessions x 2 hours = 4 hours total, across 2 days **Audience:** Employees newly issued ChatGPT licenses, first exposure to structured AI training **v4 change:** Competency framework replaced with the AI Capability Framework (Augmentation / Automation / Agentic) from AI Enablement Academy [web:71]

Learning Outcomes

By the end of this 4-hour program, participants will be able to:

1. Explain in plain language what AI/LLMs are, how they generate responses, and why confident output isn't always correct.
 2. Locate themselves on the AI Capability Framework and understand the three distinct layers of AI capability, not just their own current depth.
 3. Identify which of their own recurring work tasks are realistic candidates for AI assistance.
 4. Apply data protection rules correctly — knowing what can and cannot be entered into ChatGPT.
 5. Understand tokens and cost drivers well enough to scope prompts efficiently.
 6. Write effective prompts using a repeatable structure (Context + Task + Format), rather than trial and error.
 7. Apply a verification habit to every AI output before using it in real work.
 8. Leave with at least one live workflow already built and one committed action for the coming week.
-

What This Package Includes

Deliverable	What it gives you
2-session course plan (this document)	Full minute-by-minute agenda for both days, facilitation notes, and talking points
Slide deck outline	Section-by-section slide content ready to build out in PowerPoint/Google Slides
Prompt template library	Ready-to-use prompts participants can copy into ChatGPT on day one and start getting value from immediately
AI Capability Framework	A 3x4 grid (layer x level) used for self-positioning at the start and end of training
Quick-reference handouts	One-page cheat sheets on data protection, tokens/cost, prompting technique, and output caution
Curated learning links	Vetted, free resources for participants who want to keep building skills after the session
Facilitator checklist	Pre-session prep list so nothing is missed on the day

The intent is that a participant walks out of Day 2 with working prompt templates already applied to their real tasks — not just notes from a lecture.

Design Rationale

- Two 2-hour blocks give enough contiguous time for foundational concepts and hands-on practice without losing momentum.
- Sequence: concept → data protection → use cases + cost awareness → map to their own work → live practical examples with prompting technique → guided practice → action plan.
- The Capability Framework opens Day 1 deliberately, before any content, so participants discover — rather than are told — where they sit and how much territory exists beyond it.
- Data protection is placed right after the core concept section, before anyone touches the tool live.
- Cost/token awareness is folded into the use-case mapping exercise.
- Caution on AI outputs is introduced at the end of Day 1 and reinforced throughout Day 2's hands-on work.

DAY 1 — Foundations, Protection & Use Case Mapping (2 hours)

Time	Segment	Details
0:00–0:15	Capability Framework self-positioning	Show the full 3x4 grid. Participants locate themselves only within the Augmentation column (their license is ChatGPT, a personal-productivity tool) — almost everyone lands at L1 Foundations. Facilitator points out the Automation and Agentic columns exist entirely untouched, to make the scale of the field visible.
0:15–0:40	What AI/LLMs actually are	Plain-language coverage: pattern prediction vs. "thinking," what a prompt is, why confident-sounding output isn't always correct.
0:40–1:00	Data protection: the ground rules	Core rule: treat every AI chat like a public forum. Cover the company's licensed ChatGPT tier and its default data/training settings, never entering client data/financials/credentials, anonymizing sensitive content, and where to check the company's AI policy.
1:00–1:20	Live use-case demo (before/after)	Facilitator demonstrates 2–3 tasks the old way, then with ChatGPT — close to attendees' actual jobs.
1:20–1:45	Workflow-mapping + cost-awareness exercise (small groups)	Groups list recurring weekly tasks, flag AI candidates, and learn that longer inputs/outputs and longer conversations use more resources — scoping a task tightly is both cheaper and better.
1:45–2:00	Caution on outputs + close	Hallucination risk explained plainly. Confident tone ≠ accurate. Rule: no AI output goes out the door unchecked. Preview Day 2.

DAY 2 — Prompting, Practical Examples

& Guided Practice (2 hours)

Time	Segment	Details
0:00–0:15	Recap + framework re-check	Re-examine the grid: has anyone's Augmentation depth moved from L1 toward L2 (structured prompting with context control)? Frame L2 as this week's realistic target, not the whole grid.
0:15–0:30	How to prompt for better results	Teach Context + Task + Format formula. Emphasize specificity, breaking big asks into steps, and stating experience level so AI calibrates its answer.
0:30–1:05	Live practical examples on mapped tasks	Facilitator runs live prompts against tasks identified Day 1. Show iteration: weak prompt → refined prompt → good output. Note where a tighter prompt also reduces token usage.
1:05–1:45	Hands-on practice exercises	Participants apply ChatGPT to their own real task using the prompt templates and the Context+Task+Format formula, building a reusable prompt they keep. Facilitator reinforces "verify before you use it" on every output.
1:45–2:00	Share-outs + personal action plan	Each participant commits to one task they'll do with AI this week, plus one thing they'll always check before trusting an answer. Assign an internal "AI champion" per team.

Optional follow-up (outside the 4 hours): 15-minute "what did you try" check-in two weeks later.

The AI Capability Framework

Sourced from AI Enablement Academy's model of AI capability building [web:71]. The framework has two axes:

- **Enablement layer** (the horizontal axis) — measures the *reach* of the work, its "blast radius": **Augmentation** (personal productivity — AI helps you work),

Automation (process layer — AI runs a defined process), **Agentic** (system innovation — AI operates within a larger system).

- **Capability level** (the vertical axis) — measures how *deep* the skill goes within each layer: **L1 Foundations**, **L2 Essentials**, **L3 Fluency**, **L4 Native**.

Layer → / Level ↓	Augmentation (Personal productivity)	Automation (Process layer)	Agentic (System innovation)
L1 Foundations	Identify appropriate AI use cases in daily work and select AI tools that match task requirements.	Map existing processes to identify automation opportunities with conditional logic and error handling.	Apply agentic system design principles to architect AI-native solutions with context management and memory.
L2 Essentials	Produce AI assets and artifacts using structured prompting with context control.	Build AI-enabled interactive tools using no-code platforms.	Design organizational knowledge systems with MCP integrations, custom skills, and documented decision logic.
L3 Fluency	Design and create custom AI assistants and tools integrated into your stack.	Create multi-step workflows with agentic nodes and platform-integrated integrations.	Build production-ready agentic systems with PRD-driven development, testing protocols, and deployment.
L4 Native	Operate with an AI-first personal workflow where automation and agentic delegation are default modes.	Deploy production-grade automation infrastructure with versioning and handoff protocols.	Scale agentic systems across teams with governance frameworks and organizational memory architecture.

How to use this in the training

This 3x4 grid is a stronger "how little you know" tool than a single-track competency ladder, because it shows participants there are **three entirely separate dimensions of capability**, not just one deepening skill. A newly licensed ChatGPT user is almost always sitting at L1 Foundations in the Augmentation column only — meaning eight of the twelve cells in the grid (all of Automation, all of Agentic) are territory they haven't even entered yet.

For this specific 4-hour training:

- Only position participants within the **Augmentation** column — Automation and Agentic require tools and time this session doesn't cover. Show the full grid so the other two columns are visible, but explicitly frame them as "not today's focus."
- Realistic outcome for this training: move participants from **L1 Foundations** toward the start of **L2 Essentials** in Augmentation — i.e., from "I've tried a chatbot" to "I can produce a usable business artifact with a structured prompt."
- Use the empty Automation/Agentic columns as the explicit "here's what's still ahead" moment — this is what replaces the old five-level self-assessment as your humility hook, and it's more concrete because it names actual future skills (no-code workflow tools, agentic systems) rather than abstract "expert" labels.

Facilitation tip: Don't let the Automation/Agentic columns become a distraction or a source of anxiety — frame them as "the next chapters," not something participants are behind on. The point of showing the whole grid is scale and direction, not comparison.

Quick Reference: Data Protection Rules

- Treat every AI conversation like a public forum — nothing goes in that you wouldn't put in a press release.
 - Confirm which ChatGPT tier your company licenses (Team/Enterprise have training off by default and a data processing agreement; free/Plus consumer accounts do not).
 - Anonymize before pasting: strip names, account numbers, identifiers.
 - Never enter: client data, employee records, financials, credentials — unless company policy explicitly permits it for that tool tier.
 - Know that admin-controlled retention settings exist; data isn't kept forever, but isn't instantly erased either.
-

Quick Reference: Tokens & Cost, in Plain English

- A token is roughly $\frac{3}{4}$ of a word — about 1,000 tokens per 750 words of text.
 - Every prompt costs on both ends: input (what you type) and output (what the AI writes back) — output is typically priced several times higher.
 - Longer conversations cost more because the AI reprocesses the full history each time.
 - Clear, concise, well-scoped prompts are both cheaper and tend to produce better answers.
-

Quick Reference: How to Prompt Well

Context + Task + Format formula:

1. **Context** — who you are, the situation, relevant background.
2. **Task** — the specific thing you want done.
3. **Format** — how you want the output structured (bullets, table, tone, length).

Additional habits: be specific not vague; break large requests into steps; state your experience level; treat the first response as a draft and iterate.

Quick Reference: Caution on AI Outputs

- AI predicts plausible-sounding text; it does not verify truth.
 - Confident tone is not a signal of correctness.
 - Always human-check numbers, facts, quotes, and anything client-facing.
 - Cross-check important claims against a second, independent source.
 - Rule for the room: no AI output goes out the door unchecked.
-

Curated Learning Links (for participants who want to go further)

- **OpenAI Business Learning Hub** — official guides, real examples, and ready-to-use prompts for businesses: openai.com/business/learn
 - **OpenAI Help Center** — workspace setup, billing, security, and day-to-day usage guides for Business/Enterprise accounts: help.openai.com
 - **Learn Prompting** — the most-cited free, open-source course on prompt engineering: learnprompting.org
 - **OpenAI Prompt Engineering Guide** — official technique reference: platform.openai.com/docs/guides/prompt-engineering
 - **AI Enablement Academy** — source of the Capability Framework used in this training; useful if the organization wants to progress teams into the Automation or Agentic layers later: aienablement.academy
-

Facilitator Checklist

- ☐ Confirm ChatGPT license tier and its default data/training settings before Day 1
- ☐ Prepare 2–3 real, approved-data demo tasks relevant to this specific team
- ☐ Print or share the AI Capability Framework grid for self-positioning
- ☐ Have the company's data-privacy / acceptable-use policy ready to reference
- ☐ Prepare a "bad prompt vs. good prompt" example pair for the prompting segment
- ☐ Prepare one example of an AI hallucination to show live to make the caution segment concrete
- ☐ Print the prompt template library as handouts for Day 2
- ☐ Identify at least one enthusiastic participant to become the team's AI champion
- ☐ Schedule the optional 2-week follow-up check-in before the session ends